

# Fourier, filtering, smoothing, and noise

Nuno Vasconcelos

*ECE Department, UCSD*

*(with thanks to David Forsyth)*

# Plan for today

- ▶ we have talked about some theoretic of 2D DSP
- ▶ today we will see examples with real images
- ▶ we will:
  - talk about the discrete-space Fourier transform
  - sampling in 2D
  - visualize concepts such as frequency spectra and aliasing
  - see that very important vision operations are really just filtering, with some post processing
  - introduce some important filters that are widely used in practice
  - talk about the impact of noise and how to deal with it

# The Discrete-Space Fourier Transform

- is, once again, a straightforward extension of the 1D Discrete-Time Fourier Transform

$$X(\omega_1, \omega_2) = \sum_{n_1} \sum_{n_2} x[n_1, n_2] e^{-j\omega_1 n_1} e^{-j\omega_2 n_2}$$
$$x[n_1, n_2] = \frac{1}{(2\pi)^2} \iint X(\omega_1, \omega_2) e^{j\omega_1 n_1} e^{j\omega_2 n_2} d\omega_1 d\omega_2$$





- properties:

- basically the same as in 1D (see table in Lim, page 25)
- only novelty is separability (homework)

$$x[n_1, n_2] = x_1[n_1]x_2[n_2] \leftrightarrow X(\omega_1, \omega_2) = X_1(\omega_1)X_2(\omega_2)$$

# Image spectrum

- two images, the magnitude, and phase of their FTs



# Phase and Magnitude

## ► curious fact

- all natural images have about the same magnitude transform
- monotonically decaying with frequency


$$X(\omega_1, \omega_2) \propto \frac{1}{\omega_1^2 + \omega_2^2}$$

- hence, phase seems to matter, but magnitude largely doesn't
- we have seen this a little bit in the first problem set
- took two pictures, swap the phase transforms, compute the inverse
- here is another example




# The importance of phase

Reconstruction with **zebra**  
phase, cheetah magnitude





Reconstruction with **cheetah**  
phase, zebra magnitude



# Sampling in 2D

- ▶ consider an analog signal  $x_c(t_1, t_2)$  and let its analog Fourier transform be  $X_c(\Omega_1, \Omega_2)$ 
  - we use capital  $\Omega$  to emphasize that this is analog frequency
- ▶ sample with period  $(T_1, T_2)$  to obtain a discrete-space signal

$$x[n_1, n_2] = x_c(t_1, t_2) \Big|_{t_1=n_1T_1; t_2=n_2T_2}$$



# Sampling in 2D

- relationship between the Discrete-Space FT of  $x[n_1, n_2]$  and the FT of  $x_c(t_1, t_2)$  is simple extension of 1D result

$$X(\omega_1, \omega_2) = \frac{1}{T_1 T_2} \sum_{r_1=-\infty}^{\infty} \sum_{r_2=-\infty}^{\infty} X_c \left( \frac{\omega_1 - 2\pi r_1}{T_1}, \frac{\omega_2 - 2\pi r_2}{T_2} \right)$$


DSFT of  $x[n_1, n_2]$   
“discrete spectrum”

FT of  $x_c(\omega_1, \omega_2)$   
“analog spectrum”

- Discrete Space spectrum is sum of replicas of analog spectrum
  - in the “base replica” the analog frequency  $\Omega_1$  ( $\Omega_2$ ) is mapped into the digital frequency  $\Omega_1 T_1$  ( $\Omega_2 T_2$ )
  - discrete spectrum has periodicity  $(2\pi, 2\pi)$



# For example




$$\Omega' \rightarrow \alpha = \Omega' T_1$$

$$\Omega'' \rightarrow \beta = \Omega'' T_2$$

► no aliasing if


$$\begin{cases} \Omega' T_1 \leq 2\pi - \Omega' T_1 \\ \Omega'' T_2 \leq 2\pi - \Omega' T_2 \end{cases} \Leftrightarrow$$

$$\Leftrightarrow \boxed{\begin{cases} T_1 \leq \pi / \Omega' \\ T_2 \leq \pi / \Omega'' \end{cases}}$$



# Aliasing

- ▶ the frequency  $(\Omega'/\pi, \Omega''/\pi)$  is the critical sampling frequency
- ▶ below it we have aliasing
- ▶ this is just like the 1D case, but now there are more possibilities for overlap



# Reconstruction

- ▶ if there is no aliasing we can recover the signal in a way similar to the 1D case

$$y_c(t_1, t_2) = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} x[n_1, n_2] \frac{\sin \frac{\pi}{T_1}(t_1 - n_1 T_1)}{\frac{\pi}{T_1}(t_1 - n_1 T_1)} \frac{\sin \frac{\pi}{T_2}(t_2 - n_2 T_2)}{\frac{\pi}{T_2}(t_2 - n_2 T_2)}$$


- ▶ note: in 2D there are many more possibilities than in 1D
  - e.g. the sampling grid does not have to be rectangular, e.g. hexagonal sampling when  $T_2 = T_1/\sqrt{3}$  and

$$x[n_1, n_2] = \begin{cases} x_c(t_1, t_2)|_{t_1=n_1 T_1; t_2=n_2 T_2} & n_1, n_2 \text{ both even or odd} \\ 0 & \text{otherwise} \end{cases}$$





- in practice, however, one usually adopts the rectangular grid

# Sampling and aliasing in 2D

- ▶ in summary, sampling is not very different from the 1D case




- ▶ but aliasing is a lot more fun to look at in images
- ▶ it shows up in video too (the wagon wheel effect)



- ▶ a sequence of images obtained by down-sampling **without any filtering**
- ▶ aliasing: the **low-frequency parts are replicated throughout the low-res image**



# The role of smoothing



- ▶ too little leads to **aliasing**
- ▶ too much leads to **loss of information**

# Linear Filtering

- ▶ smoothing is implemented with linear filters
- ▶ given an image  $x(n_1, n_2)$ , filtering is the process of convolving it with a kernel  $h(n_1, n_2)$

$$y(n_1, n_2) = \sum_{k_1 k_2} x(k_1, k_2) h(n_1 - k_1, n_2 - k_2)$$


- ▶ some very common operations in image processing are nothing but filtering, e.g.
  - smoothing an image by low-pass filtering
  - contrast enhancement by high pass filtering
  - finding image derivatives
  - noise reduction

# Popular filters


## ► box function

$$R_{N_1 \times N_2}(n_1, n_2) = \begin{cases} 1, & 0 \leq n_1 \leq N_1 - 1, 0 \leq n_2 \leq N_2 - 1 \\ 0 & \text{otherwise} \end{cases}$$

## ► Fourier transform of a box is the **sinc**, low-pass filter




$\mathcal{F}$




## ► side-lobes produce artifacts, smoothed image does not look like the result of defocusing

# Example: Smoothing by Averaging



# Camera defocusing

- ▶ if you point an out-of-focus camera at a very small white light (e.g. a light-bulb) at night, you get something like this
- ▶ the light can be thought of as an **impulse**
- ▶ this must be the **impulse response**
- ▶ well approximated by a **Gaussian**
- ▶ more natural filter for image blur than the box



$$h(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right)$$



# The Gaussian

- ▶ the discrete space version is

$$h(n_1, n_2) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{n_1^2 + n_2^2}{2\sigma^2}\right)$$

- ▶ obviously separable


$$h(n_1, n_2) = \underbrace{\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{n_1^2}{2\sigma^2}}}_{h(n_1)} \times \underbrace{\frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{n_2^2}{2\sigma^2}}}_{h(n_2)}$$

- ▶  $h(n_1, n_2)$  has Fourier transform

$$H(\varpi_1, \varpi_2) = \exp\left(-\frac{\sigma^2(\varpi_1^2 + \varpi_2^2)}{2}\right)$$


# The Gaussian filter

- the Fourier transform of a Gaussian is a Gaussian  
 $(\sigma_x, \sigma_y) \propto (1/\sigma_{w1}, 1/\sigma_{w2})$




- note that there are **no annoying side-lobes**

# Smoothing with a Gaussian



# Role of the variance




- ▶ the variance controls the amount of smoothing
- ▶ each column shows different realizations of an image of gaussian noise
- ▶ each row shows smoothing with gaussians of different  $\sigma$

# Gradients and edges

- ▶ for image understanding, one of the problems is that there is **too much information** in an image
- ▶ just smoothing is not good enough
- ▶ how to **detect important (most informative) image points**?
- ▶ note that **derivatives are large at points of great change**
  - changes in reflectance (e.g. checkerboard pattern)
  - change in object (an object boundary is different from background)
  - change in illumination (the boundary of a shadow)
- ▶ these are usually called **edge points**
- ▶ detecting them could be useful for various problems
  - segmentation: we want to know what are **object boundaries**
  - recognition: **cartoons are easy to recognize** and terribly efficient to transmit



# The importance of edges



# Gradients

- ▶ for a 2D function,  $f(x,y)$  the **gradient** at a point  $(x_0, y_0)$

$$\begin{aligned}\nabla f(x_0, y_0) &= \left( \frac{\partial f}{\partial x}(x_0, y_0), \frac{\partial f}{\partial y}(x_0, y_0) \right)^T \\ &= \left( f_x(x_0, y_0), f_y(x_0, y_0) \right)^T\end{aligned}$$



is the **direction of greatest increase** at that point

- ▶ the **gradient magnitude**

$$\|\nabla f(x_0, y_0)\|^2 = \left( \frac{\partial f}{\partial x}(x_0, y_0) \right)^2 + \left( \frac{\partial f}{\partial y}(x_0, y_0) \right)^2$$

**measures the rate of change**

- ▶ it is **large at edges!**



- large gradient magnitude
- small gradient magnitude

# Derivatives and convolution

- recall that a **derivative** is defined as

$$\frac{\partial f(x)}{\partial x} = \lim_{\Delta x \rightarrow 0} \frac{f(x + \Delta x) - f(x)}{\Delta x}$$

- linear and shift invariant, so must be the result of a convolution.
- we could **approximate** as

$$\frac{\partial f(n)}{\partial n} = \frac{f(n+1) - f(n)}{1} = f(n+1) - f(n) = f * h(n)$$


- where the **derivative kernel** is

$$h(n) = \delta(n+1) - \delta(n)$$

# Finite difference kernels


- ▶ in two dimensions we have various possible kernels
- ▶ e.g. ,  $N_1=2$ ,  $N_2=3$ , derivative **along  $n_1$** , (line  $n_2=k$ ) (horizontal)

$$\begin{array}{cc} 0 & 0 \\ 1 & -1 \\ 0 & 0 \end{array} \quad \begin{array}{cc} 1 & -1 \\ 1 & -1 \\ 1 & -1 \end{array}$$




- ▶ derivative **along  $n_2$** , (line  $n_1=k$ ) (vertical)

$$\begin{array}{ccc} 0 & -1 & 0 \\ 0 & 1 & 0 \end{array} \quad \begin{array}{ccc} -1 & -1 & -1 \\ 1 & 1 & 1 \end{array}$$



- ▶ derivative **along line  $n_1=n_2$**  (diagonal)

$$\begin{array}{ccc} 0 & 0 & -1 \\ 0 & 0 & 0 \\ 1 & 0 & 0 \end{array} \quad \begin{array}{ccc} 1 & -1 & 0 \\ 0 & 1 & -1 \\ 0 & 0 & 1 \end{array}$$



# Finite difference kernels


- note that, when

$$h(n_1, n_2) = \begin{bmatrix} 0 & 0 \\ 1 & -1 \\ 0 & 0 \end{bmatrix}$$

- we have

$$\begin{aligned} H(\omega_1, \omega_2) &= e^{j\omega_1} - 1 \\ &= \left( e^{j\frac{\omega_1}{2}} - e^{-j\frac{\omega_1}{2}} \right) e^{j\frac{\omega_1}{2}} \\ &= 2 \sin\left(\frac{\omega_1}{2}\right) e^{j\frac{\omega_1}{2}} \end{aligned}$$


- derivative is a **high-pass filter**
- hw: check that **this holds for all others**
- intuitive, because a derivative is a measure of the **rate of change of a function**





# Finite differences

► Q: which one do we have here? (gray=0, white=+, dark=-)



# Finite differences and noise

- ▶ because they perform high-pass filtering, finite difference filters respond strongly to noise
- ▶ generally, the larger the noise the stronger the response
- ▶ for noisy images it is usually best to apply some smoothing before computing derivatives
- ▶ what do mean by noise?
- ▶ we only consider the simplest model
  - independent stationary additive Gaussian noise
  - the noise value at each pixel is given by an independent draw from the same normal probability distribution

$$y(n_1, n_2) = x(n_1, n_2) + \varepsilon(n_1, n_2), \quad \varepsilon \sim \mathcal{N}(0, \sigma^2)$$

$\sigma=1$




$\sigma=16$



# Finite differences responding to noise

increasing noise variance



- note that as the noise variance increases the estimates of the image derivative are also very noisy
- Q: would a larger filter do better?

# The response of a linear filter to noise

- ▶ suppose we apply filter  $h(n_1, n_2)$  to

$$y(n_1, n_2) = x(n_1, n_2) + \varepsilon(n_1, n_2), \quad \varepsilon \sim N(0, \sigma^2)$$

- ▶ by linearity

$$h * y = \underbrace{h * x}_{\text{output signal}} + \underbrace{h * \varepsilon}_{\text{output noise}}$$

- ▶ consider the noise term

$$v(n_1, n_2) = \sum_{k_1 k_2} \varepsilon(k_1, k_2) h(n_1 - k_1, n_2 - k_2)$$

- ▶ by linearity of expectation

$$E[v(n_1, n_2)] = \sum_{k_1 k_2} E[\varepsilon(k_1, k_2)] h(n_1 - k_1, n_2 - k_2) = 0$$

# The response of a linear filter to noise

► and since the  $\varepsilon(n_1, n_2)$  are independent

$$\begin{aligned}\text{var}[v(n_1, n_2)] &= \sum_{k_1 k_2} \text{var}[\varepsilon(k_1, k_2)] h^2(n_1 - k_1, n_2 - k_2) \\ &= \sigma^2 \sum_{k_1 k_2} h^2(n_1 - k_1, n_2 - k_2)\end{aligned}$$

► output error has

- zero mean
- variance that increases with the size of the filter and the input variance

► increasing the filter size would definitely not help!

# Smoothing reduces noise

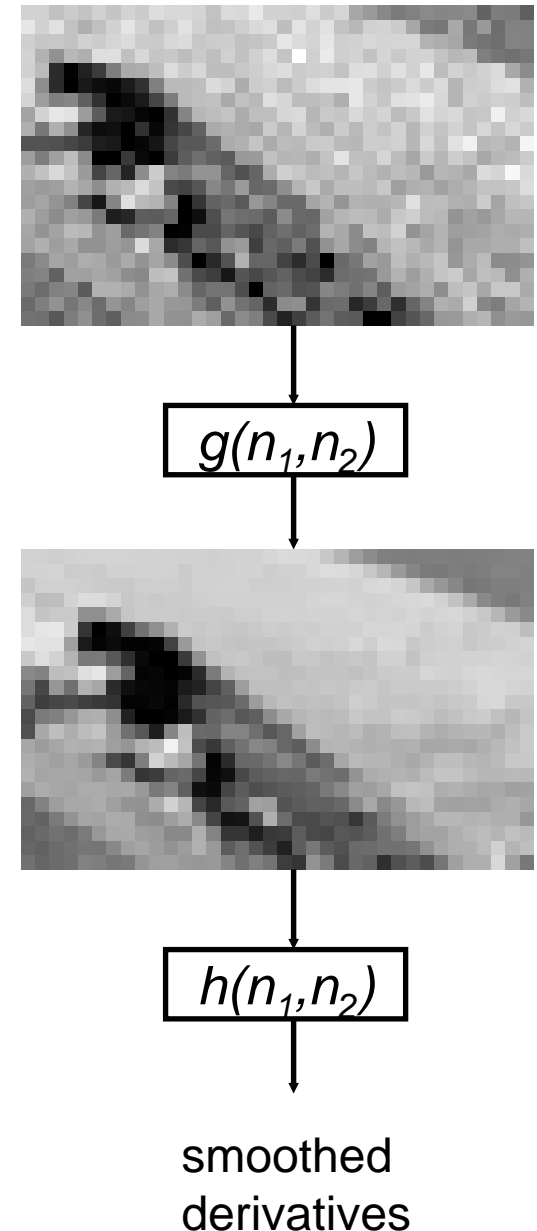
- ▶ noise has a lot of high-frequencies
- ▶ strategy:
  1. start by low-pass filtering, to suppress noise
  2. compute derivative on smoothed image
- ▶ i.e. for a smoothing filter  $g(n_1, n_2)$  compute

$$h * (g * x)$$

- ▶ note that, by associativity of convolution, this is equal to

$$(h * g) * x$$

- ▶ i.e. filter the image with the filter  $h * g$





# The derivative of a Gaussian

- ▶ let's consider, for example,

$$h(n_1, n_2) = \delta(n_1 + 1, n_2) - \delta(n_1, n_2)$$


- ▶ in which case

$$h * g(n_1, n_2) = g(n_1 + 1, n_2) - g(n_1, n_2)$$


is a **difference of two Gaussians**

- ▶ this is the **derivative of a Gaussian (DoG) filter**
- ▶ for other definitions of  $h$  we have a similar result

DoG along  $n_1$




DoG along  $n_2$




# Smoothed derivatives

no smoothing




with smoothing




# Choosing the right scale

- ▶ the scale of the smoothing filter affects
  - derivative estimates,
  - the semantics of the derivative image
- ▶ trade-off between noise and ability to detect detail


1 pixel



3 pixels




7 pixels



**Any questions?**

# Gradients and edges

- ▶ by now we have a good idea of how to differentiate images
- ▶ remember that **edges are points of large gradient magnitude**
- ▶ **edge detection strategy**
  1. determine **magnitude of image gradient**
$$\|\nabla f(x_0, y_0)\|^2 = \left(\frac{\partial f}{\partial x}(x_0, y_0)\right)^2 + \left(\frac{\partial f}{\partial y}(x_0, y_0)\right)^2$$
  2. mark **points where gradient magnitude is particularly large wrt neighbours** (ideally, curves of such points)



- large gradient magnitude
- small gradient magnitude

# Looks easy but

- ▶ three major issues (to discuss next class):
  - 1) gradient magnitude at different scales is different (see below); which should we choose?
  - 2) gradient magnitude is large along thick trail; what are the significant points?
  - 3) how do we link the relevant points up into curves?

